

CARM: Controlled Attention Routing and Masking

Architectural Concept Isolation in Hybrid AI Systems

Horacio López Barrios

Independent Researcher

`h@ual.li`

2026 *Preprint — work in progress*

Abstract

We introduce **Controlled Attention Routing and Masking (CARM)**, a mechanism for enforcing concept-level access policies in transformer attention at the architectural layer rather than through behavioural training or prompt engineering. CARM operates by maintaining a routing table indexed on semantic concept identifiers — in the present implementation, 32-bit hierarchical identifiers following the TOSID scheme — and applying the resulting access mask prior to attention score computation, so that blocked concepts receive exactly zero attention weight in every head, in every pass, regardless of query content.

We situate CARM within a two-component AI architecture comprising an **ACI** (Artificial Contextual Intelligence) layer — a non-deterministic reasoning component, typically a large language model — and an **AXI** (Artificial eXpert Intelligence) layer — a deterministic, domain-compiled expert component. CARM defines the controlled interface between these layers: AXI supplies the routing policy; ACI operates within the concept space that policy permits.

We present a proof-of-concept implementation in C demonstrating: (i) correct routing enforcement across 5,000 independent attention computations with zero violations; (ii) measured routing latency of 31–159 ns across concept counts from 5 to 100, against attention computation times of 25–366 μ s over the same range — noting that the lower end of the routing latency range approaches system timer resolution and should be interpreted as an order-of-magnitude bound rather than a precise figure; at worst case the routing overhead is below 0.7% of total computation time; (iii) latency distribution stability with p_{99} of 54 ns across 20,000 samples; and (iv) monotonically increasing output divergence as blocked concept fraction increases, confirming that policy changes produce semantically meaningful effects on the output embedding.

This work establishes CARM as a formally defined, measurable mechanism and demonstrates its viability at the attention-component level. Integration with complete transformer architectures — full encoder/decoder stacks, autoregressive generation, production-scale vocabularies — is ongoing and constitutes the primary open direction of current research.

Contents

1	Introduction	3
1.1	The constraint problem in current AI	3
1.2	Architectural versus behavioural enforcement	3
1.3	The ACI/AXI architecture	3
1.4	TOSID as semantic identifier	4
1.5	Scope and contributions	4
2	The CARM Mechanism	4
2.1	Informal statement	4
2.2	Formal definition	5
2.3	Discussion of the criteria	6
2.4	What CARM does not specify	6
2.5	Relationship to attention masking in existing systems	7
3	Architecture	7
3.1	Component roles	7
3.2	The CARM interface	7
3.3	Policy representation	8
3.4	Deployment contexts	8
3.5	What this architecture does not yet address	9
4	Implementation	9
4.1	Overview	9
4.2	The routing layer	9
4.2.1	TOSID structure	9
4.2.2	The routing table	10
4.2.3	Context switching	10
4.3	The attention layer	10
4.3.1	Configuration	10
4.3.2	The three-phase computation	11
4.4	Concept embeddings	12
4.4.1	Phase 1.1: synthetic clustered embeddings	12
4.4.2	Phase 1.2: GloVe embeddings at scale	12
4.5	Test concept database	12
4.6	Known issues and pre-publication fixes required	13
5	Results	13
5.1	Experimental setup	13
5.2	Correctness: zero information leakage	13
5.3	Performance: routing overhead	14
5.4	Latency stability	14
5.5	Policy effect on output	15
5.6	Summary	16
6	Discussion and Ongoing Work	16
6.1	What has been established	16
6.2	Identifier isolation and semantic isolation	17
6.3	Limitations of the current implementation	17
6.4	The feedback path and TOSIDcoders	18
6.5	The research programme	19
6.6	Relation to broader AI safety work	20
6.7	Conclusion	20

1 Introduction

1.1 The constraint problem in current AI

Contemporary large language model (LLM) deployment relies heavily on behavioural mechanisms for concept restriction: system prompts instruct the model to avoid certain topics; reinforcement learning from human feedback (RLHF) [2] shapes the probability distribution away from undesired outputs; constitutional AI [3] embeds refusal behaviour during fine-tuning. These approaches share a structural limitation: they operate on the *output* of the model’s reasoning process, or on its training distribution, rather than on the reasoning process itself. A sufficiently adversarial prompt, or a distribution shift, can move the model back into suppressed territory. The constraint is behavioural — a tendency, not a guarantee.

The problem is not that these techniques are poorly engineered. The problem is that they are fighting the fundamental nature of a stochastic system. A dense neural network trained on a large corpus develops entangled representations in which concepts are not cleanly separable. Suppressing one concept through training nudges probabilities; it does not excise the concept from the representational space.

This has practical consequences in domains where concept access must be governed by policy rather than tendency: a clinical decision-support system must not leak Protected Health Information (PHI) into outputs accessible to billing staff; a financial AI operating under information-barrier obligations must not attend to material non-public information when generating advice for retail clients; a legal AI must not cite case law from an inapplicable jurisdiction. In each case, the required guarantee is categorical, not probabilistic.

1.2 Architectural versus behavioural enforcement

We distinguish two modes of concept restriction:

Behavioural enforcement operates on model weights or outputs. It includes prompt engineering, RLHF, fine-tuning, and output filtering. It is flexible, easy to deploy on existing models, and insufficient for categorical guarantees.

Architectural enforcement operates on the computation graph. It constrains what the model *can* attend to, not what it *tends* to attend to. A concept that receives zero attention weight contributes nothing to the output embedding, regardless of what the query contains, regardless of how the prompt is phrased, regardless of the model’s training distribution.

CARM is an architectural enforcement mechanism. It does not modify model weights. It does not require retraining. It interposes a routing layer between the concept vocabulary and the attention computation, ensuring that blocked concepts are excluded from the softmax domain before any scoring occurs.

1.3 The ACI/AXI architecture

We motivate CARM within a hybrid AI architecture that we term the ACI/AXI model.

An **ACI (Artificial Contextual Intelligence)** component provides flexible, context-sensitive reasoning — the kind of capability that LLMs excel at: language understanding, situational inference, handling of novel inputs, generation of coherent text. ACI components are inherently non-deterministic; their value lies precisely in their ability to generalise beyond what was explicitly programmed.

An **AXI (Artificial eXpert Intelligence)** component provides deterministic, verifiable, domain-compiled expertise — analogous in spirit to classical expert systems [6], but without their brittleness. An AXI component encodes domain knowledge in a form that can be queried reliably, audited completely, and updated without retraining. AXI components do not generalise; they enforce.

In the ACI/AXI model, the two components are complementary rather than competing. ACI handles the parts of a task that require contextual flexibility; AXI handles the parts that require precision and policy compliance. CARM defines the interface at which AXI policy constrains ACI attention: the routing table is an AXI artefact; the attention computation is an ACI operation; CARM ensures the former governs the latter.

This paper does not present a full ACI/AXI system. It presents the CARM mechanism that makes such a system possible, and demonstrates its correctness and performance properties at the component level.

1.4 TOSID as semantic identifier

CARM requires a scheme for identifying concepts. Any scheme that assigns a stable, prefix-matchable identifier to each concept in the vocabulary is compatible with CARM. In the present implementation we use a simplified 32-bit form of the Taxonomic Ontological Semantic IDentification (TOSID) scheme, in which the upper 16 bits encode domain and category — enabling $O(1)$ prefix-matched lookup in a 64 KB routing array — and the lower 16 bits encode the specific concept variant.

The full TOSID specification is a separately developed system and is not published here. The present paper requires only the structural properties that TOSID provides: hierarchical organisation, prefix-matchability, and stable assignment. Any identifier system with these properties is compatible with CARM.

1.5 Scope and contributions

This paper makes the following contributions:

1. **A formal definition of CARM** as a mechanism, with stated criteria for what constitutes a CARM-compliant implementation (section 2).
2. **Introduction of the ACI/AXI architectural distinction** as the context in which CARM operates, with minimal but sufficient definitions (section 3).
3. **A proof-of-concept implementation** in C demonstrating CARM operating over a multi-head attention layer with real semantic embeddings (section 4).
4. **Empirical results** establishing correctness (zero leakage across 5,000 independent measurements), performance (routing latency 31–159 ns against attention computation times of 25–366 μ s; worst-case overhead below 0.7%), and latency stability (p_{99} of 54 ns across 20,000 samples) — with explicit acknowledgement that the lower latency measurements approach timer resolution and are reported as order-of-magnitude bounds (section 5).
5. **An account of ongoing and planned work**, including the path to integration with complete transformer stacks (section 6).

We make no claim that CARM is sufficient, as implemented here, for production deployment with a full LLM. We claim that it is a well-defined, correctly operating mechanism at the component level, and that the path from here to full integration is clear and tractable.

The implementation is available at: <https://github.com/ha1tch/carm>

2 The CARM Mechanism

2.1 Informal statement

CARM is a mechanism that interposes a policy-controlled access filter between a concept vocabulary and the attention computation that consumes it. The filter is evaluated once per

attention pass, prior to any score computation, and produces a binary decision — allow or deny — for each concept in the vocabulary. Denied concepts are excluded from the softmax domain entirely: they are not scored, they receive no attention weight, and they contribute nothing to the output embedding. The filter is supplied by an external policy authority; the attention computation has no knowledge of, and no influence over, the filter’s decisions.

2.2 Formal definition

Let $\mathcal{V} = \{v_1, v_2, \dots, v_n\}$ be a concept vocabulary, where each concept v_i is associated with a stable semantic identifier $\text{id}(v_i) \in \mathcal{I}$ for some identifier space $\mathcal{I}$.

Let $P : \mathcal{I} \rightarrow \{\text{ALLOW}, \text{DENY}\}$ be a *routing policy*, a total function from identifiers to access decisions. P is supplied by an external authority and is defined independently of any query or attention computation.

Let q be a query embedding. A standard scaled dot-product attention computation over $\mathcal{V}$ with query q produces attention weights:

$$\alpha_i = \text{softmax}\left(\frac{q \cdot k_i}{\sqrt{d}}\right)_i$$

where k_i is the key projection of v_i and d is the key dimension.

A **CARM-compliant** attention computation modifies this as follows:

Phase 1. Routing phase (prior to any score computation): construct the accessible set

$$\mathcal{A} = \{v_i \in \mathcal{V} : P(\text{id}(v_i)) = \text{ALLOW}\}.$$

This phase must be independent of q : the composition of $\mathcal{A}$ must not depend on the query embedding, the attention weights, or any output of the current or prior attention pass.

Phase 2. Attention phase: compute attention scores and weights over $\mathcal{A}$ only:

$$\alpha_i = \begin{cases} \text{softmax}(q \cdot k_i / \sqrt{d})_i & \text{if } v_i \in \mathcal{A}, \\ 0 & \text{if } v_i \notin \mathcal{A}. \end{cases}$$

Phase 3. Output phase: compute the output embedding as a weighted sum over $\mathcal{A}$ only:

$$\mathbf{out} = \sum_{v_i \in \mathcal{A}} \alpha_i \cdot v_i.$$

Definition 2.1 (CARM compliance). An attention implementation is *CARM-compliant* if and only if all four of the following criteria hold:

- (C1) **Zero weight.** For every $v_i \notin \mathcal{A}$, the attention weight $\alpha_i = 0$ exactly, not approximately.
- (C2) **Query independence of routing.** The construction of $\mathcal{A}$ does not depend on q or on any attention score or output.
- (C3) **Policy externality.** P is supplied by an authority external to the attention computation. The attention component does not modify, override, or circumvent P .
- (C4) **Completeness.** Every concept in $\mathcal{V}$ is evaluated by P before the attention phase begins. No concept bypasses routing.

2.3 Discussion of the criteria

C1 (Zero weight) is the condition that distinguishes architectural from behavioural enforcement. A soft masking approach — multiplying attention logits by a small value, adding a large negative bias, or clipping weights — satisfies neither C1 nor the security model. A concept with any nonzero attention weight, however small, leaks information into the output embedding. CARM requires exact zero, achieved by excluding blocked concepts from the softmax domain entirely rather than by score manipulation.

C1, in conjunction with C2–C4, yields two guarantees that characterise the relationship between the AXI policy layer and the ACI reasoning layer:

Within a CARM-compliant attention layer, C1–C4 jointly guarantee that AXI policy is the sole determinant of which concepts are accessible to ACI reasoning — a mathematical consequence of excluding denied concepts from the softmax domain prior to score computation.

CARM formally guarantees that within a compliant attention layer, no denied concept contributes to ACI output — by construction, not by tendency.

Both guarantees are scoped to the attention layer as demonstrated in this work. Whether they extend to a full transformer stack — including residual connections, feed-forward sublayers, and cross-attention paths not governed by CARM — is an open question and a primary motivation for the full-stack integration work described in section 6.

C2 (Query independence) closes the primary attack surface. If the composition of $\mathcal{A}$ could depend on the query, an adversary could craft queries that shift concepts from DENY to ALLOW. Routing must be a pure function of concept identifiers and the externally supplied policy, evaluated before any query processing occurs.

C3 (Policy externality) is what makes CARM useful in the ACI/AXI architecture. The attention component — an ACI element — cannot be trusted to self-regulate, not because it is malicious, but because it is stochastic and its behaviour under novel inputs is not fully predictable. The policy must come from outside. In the ACI/AXI model, P is an AXI artefact: deterministic, auditable, and updated independently of model training.

C4 (Completeness) prevents a class of implementation errors in which newly added concepts, or concepts outside the routing table’s defined domain, default to ALLOW. The correct default is DENY: any concept not explicitly permitted by P must be excluded from $\mathcal{A}$. This is the closed-world assumption applied to concept access.

2.4 What CARM does not specify

CARM is a mechanism specification, not an implementation specification. It does not prescribe:

- **The identifier scheme.** Any stable, prefix-matchable identifier space is compatible. The present implementation uses 32-bit TOSID identifiers; others are possible.
- **The routing table data structure.** The $O(1)$ array lookup used here is one choice. A trie, a hash map, or a hardware TCAM are equally CARM-compliant provided C1–C4 are satisfied.
- **The policy authority.** Who supplies P , how P is updated, and what audit trail P carries are system-level concerns, not CARM concerns.
- **The number of heads.** CARM applies per-head but the routing phase is evaluated once and the result shared across all heads in a multi-head attention computation.

- **The embedding space.** CARM is agnostic to whether concepts are represented by trained embeddings, handcrafted embeddings, or binary vectors, provided the attention computation is otherwise standard.

2.5 Relationship to attention masking in existing systems

Attention masking is not new. Transformer implementations routinely apply masks for causal (autoregressive) attention, padding, and cross-attention scope [1]. CARM is distinguished from these by three properties.

First, standard masks are **structural**: they encode sequence position relationships (causality, padding boundaries). CARM masks are **semantic**: they encode domain policy over concept meaning.

Second, standard masks are **static with respect to policy**: a causal mask is determined by sequence length and does not change based on who is querying or what domain they are operating in. CARM masks are **context-switched**: the same model can operate under different policies in different deployment contexts without retraining.

Third, standard masks have **no external authority requirement**: they are computed internally from sequence structure. CARM masks are supplied by an **external policy authority** (C3), establishing a clean separation between the reasoning component and the policy component.

These distinctions make CARM a new mechanism rather than a restatement of existing practice, notwithstanding the use of masking as the underlying primitive.

3 Architecture

3.1 Component roles

The ACI/AXI architecture assigns distinct roles to its two components, which are not interchangeable and are not in competition.

The **ACI component** receives a query, attends over a concept space, and produces a contextually informed output embedding. It is permitted to be stochastic, to generalise beyond its training distribution, and to produce outputs that could not have been explicitly anticipated. An ACI component is not trusted to self-restrict.

The **AXI component** maintains domain knowledge in a deterministic, auditable form and supplies policy to the system. It does not reason; it enforces. An AXI component is trusted precisely because it does not generalise: its outputs for any given input are fixed, verifiable, and independent of the ACI component’s state.

Neither component is subordinate in the general sense. Each is authoritative within its domain: ACI over contextual reasoning, AXI over policy and domain correctness. CARM is the mechanism that makes this division enforceable at the computational level.

3.2 The CARM interface

CARM sits at the boundary between the AXI and ACI components. Figure 1 shows the data flow.

Three properties of this arrangement are structurally significant.

Unidirectional policy flow. Policy flows from AXI to CARM to ACI. There is no feedback path from ACI to CARM or from ACI to AXI. The ACI component cannot influence its own concept access.

Query isolation from routing. The query q enters the system below the CARM layer. It is visible only to the ACI component and only after the accessible set $\mathcal{A}$ has been determined. This enforces C2 at the architectural level, not merely as an implementation convention.




Figure 1: The AXI/CARM/ACI architecture. Policy flows unidirectionally from AXI through the CARM layer to ACI. The dashed inhibiting arrow indicates the absence of any feedback path from ACI to the routing layer.

Single evaluation per pass. The routing phase is evaluated once per attention pass, producing $\mathcal{A}$ before any head begins score computation. All heads in a multi-head attention computation share the same accessible set. This is both an efficiency property and a consistency property: it is not possible for different heads to operate under different policies within the same pass.

3.3 Policy representation

The routing policy P is represented in the implementation as a 64KB array indexed by the upper 16 bits of each concept’s TOSID identifier. Each entry is a single byte: 0 for DENY, 1 for ALLOW. This structure is an AXI artefact: it is built from domain rules, not learned from data, and it can be inspected, versioned, and audited independently of the ACI component.

The 64KB size is not incidental. It fits within the L2 cache of contemporary processors, making repeated lookups effectively free in terms of memory latency. The routing overhead figures reported in section 5 reflect this: the dominant cost is iterating over the concept vocabulary to apply the lookup, not the lookup itself.

Policy rules are expressed at the domain-and-category level (upper 16 bits of the identifier), meaning a single rule covers all concept variants within a domain category. Finer-grained rules — down to individual concept variants — are possible within the same structure. This mirrors the hierarchical structure of network routing tables, which was a deliberate design choice: the operational model for CARM policy management is closer to firewall administration than to model fine-tuning.

3.4 Deployment contexts

A single ACI component can operate under multiple AXI policies by switching the routing table between passes. No retraining, no weight modification, no reloading of the model is required. The policy switch is a pointer replacement and a cache warm-up — an operation taking microseconds, not minutes.

This context-switching property is what makes CARM practically useful in multi-tenant or multi-role deployments. A clinical AI system serving both treating clinicians and billing staff can

apply a permissive medical policy for the former and a restricted policy blocking PHI for the latter, using the same underlying ACI component with different CARM configurations.

3.5 What this architecture does not yet address

The architecture as described operates at the level of a single attention layer over an explicit concept vocabulary. A production deployment with a full transformer-based LLM presents additional considerations not yet addressed:

- **Implicit concept representation.** In a trained LLM, concepts are not stored as discrete vocabulary entries with explicit identifiers. They are distributed across token embeddings and emerge from the interaction of layers. Applying CARM requires either a concept-to-token mapping layer, or integration at the token embedding level with TOSID-annotated vocabularies.
- **Residual and feed-forward paths.** A transformer block contains paths beyond the attention sublayer. Whether blocked concepts can re-enter the computation through residual connections or feed-forward activations is an empirical question that full-stack integration must address.
- **Autoregressive generation.** Sequential token generation introduces context accumulation dynamics not present in single-pass attention. The routing policy must remain effective as context grows, and must not be bypassable through adversarial prompt construction. This is the subject of ongoing work described in section 6.

These are not objections to the mechanism; they are the open research questions that the mechanism’s existence makes it possible to ask precisely.

4 Implementation

4.1 Overview

The proof-of-concept implementation is written in C (C11), compiled with `gcc -O2 -march=native`, and targets Linux on x86-64. It is self-contained: no machine learning framework, no GPU, no external numerical library beyond `libm`. This is a deliberate choice. The goal is to demonstrate the CARM mechanism in the clearest possible form, without the implementation being entangled with framework conventions that would obscure what is structurally necessary.

The implementation consists of two separable components:

1. **The routing layer** (`routing_table.h`, `routing_table.c`): the $O(1)$ TOSID-indexed policy array and the rule application machinery.
2. **The attention layer** (`multihead_attention.h`, `multihead_attention.c`): the multi-head scaled dot-product attention computation, modified to enforce CARM compliance.

The separation is intentional and reflects the AXI/ACI boundary: the routing layer is an AXI artefact, built from rules, inspectable and auditable; the attention layer is the ACI computation, which receives the accessible set $\mathcal{A}$ and operates only within it.

4.2 The routing layer

4.2.1 TOSID structure

The implementation uses a 32-bit simplified TOSID encoding:

```

1 // 0xDDCCVVVV
2 // DD    = Domain      (8 bits, bits 31-24)
3 // CC    = Category    (8 bits, bits 23-16)
4 // VVVV = Variant      (16 bits, bits 15-0)
5 //
6 // Routing operates on the top 16 bits (DD + CC) only.
7 // A single rule covers all variants within a domain-category pair.
8
9 typedef uint32_t TOSID;

```

Listing 1: TOSID bit layout (32-bit simplified form).

The upper 16 bits (domain + category) serve as the routing key. This means a single policy rule covers all concept variants within a domain-category pair — semantically natural, since policy is typically expressed at the level of domain and category rather than individual concept instances.

4.2.2 The routing table

The routing table is a flat 64 KB array indexed by the upper 16 bits of the TOSID. Each entry is a single byte: `ROUTE_ALLOW` (1) or `ROUTE_DENY` (0). The default on initialisation is `ROUTE_DENY` for all 65,536 entries, implementing C4 (completeness with default-deny).

```

1 #define ROUTING_TABLE_SIZE 65536 // 2^16
2
3 typedef struct {
4     uint8_t actions[ROUTING_TABLE_SIZE]; // 64 KB
5     char    context_name[64];
6     int     rule_count;
7 } FastRoutingTable;
8
9 // O(1) lookup --- inlined for maximum performance
10 static inline int routing_table_check(
11     uint32_t tosid, const FastRoutingTable* table) {
12     uint16_t index = (tosid >> 16) & 0xFFFF;
13     return table->actions[index] == ROUTE_ALLOW;
14 }

```

Listing 2: Routing table structure and O(1) lookup.

Rules are applied at construction time by iterating over all 65,536 prefix indices and marking those that match the rule’s prefix/mask pair. This is a one-time $O(n)$ build cost; subsequent lookups are $O(1)$ regardless of the number of rules.

4.2.3 Context switching

A routing context (e.g. clinical, financial, unrestricted) is a named `FastRoutingTable` instance. Switching policy between passes requires only replacing the pointer to the active table — no model reload, no retraining. The 64 KB footprint fits comfortably in L2 cache on all contemporary server processors, so warm lookups add negligible latency.

4.3 The attention layer

4.3.1 Configuration

The proof-of-concept operates with the following parameters:

Parameter	Value
Embedding dimension	64
Number of heads	8
Head dimension	8 (= 64 / 8)
Maximum concept vocabulary	100
Weight initialisation	Xavier (per-head seed for diversity)

These are deliberately small. The goal is mechanical clarity, not scale. Phase 1.2 of the implementation programme extended the vocabulary to 1,000 concepts with 100-dimensional GloVe embeddings and 10 attention heads; that work is described in section 4.4.2.

4.3.2 The three-phase computation

The core function `compute_multihead_attention_with_routing` implements the three CARM phases in strict sequence. Phase 1 completes entirely before Phase 2 begins; Phase 2 operates exclusively over the accessible set produced by Phase 1.

Phase 1 — Routing enforcement (once, before any head):

```

1 // Executed once for all heads --- before any score computation
2 Concept* accessible[MAX_CONCEPTS];
3 int nacc = 0, nblocked = 0;
4
5 for (int i = 0; i < concept_count; i++) {
6     if (check_routing_access(concepts[i].tosid, routing_table)) {
7         accessible[nacc++] = &concepts[i];
8     } else {
9         result->blocked_tosids[nblocked++] = concepts[i].tosid;
10    }
11 }
12 // accessible[] now contains only ALLOW concepts.
13 // blocked concepts are not passed to any head.
```

Listing 3: Phase 1: routing enforcement, query-independent.

The critical property here is that `check_routing_access` receives only the concept’s TOSID and the routing table. It has no access to the query embedding, the attention scores, or any prior output. C2 is enforced structurally, not by convention.

Phase 2 — Attention over accessible concepts only:

```

1 float scale = 1.0f / sqrtf((float)HEAD_DIM);
2
3 for (int h = 0; h < NUM_HEADS; h++) {
4     // Project keys and values for accessible concepts only
5     float keys[MAX_CONCEPTS][HEAD_DIM];
6     float vals[MAX_CONCEPTS][HEAD_DIM];
7     for (int i = 0; i < nacc; i++) {
8         project_embedding(accessible[i]->embedding,
9                             keys[i], W_key[h], HEAD_DIM);
10        project_embedding(accessible[i]->embedding,
11                            vals[i], W_val[h], HEAD_DIM);
12    }
13    // Attention scores over accessible concepts only
14    float scores[MAX_CONCEPTS];
15    for (int i = 0; i < nacc; i++) {
16        scores[i] = scale * dot(query_h, keys[i], HEAD_DIM);
17    }
18    softmax(scores, nacc); // softmax domain = accessible only
```

```

19 // Blocked concepts have no entry in scores[].
20 // Their attention weight is identically zero.
21 }

```

Listing 4: Phase 2: scaled dot-product attention over `accessible[]` only.

The softmax is computed over `nacc` entries — the accessible concepts — not over the full vocabulary. Blocked concepts have no entry in `scores[]` at all; they are not assigned zero, they are absent. The zero weight property (C1) follows from their absence from the softmax domain, not from any explicit zeroing step.

Phase 3 — Concatenation and output projection:

Head outputs are concatenated and projected through a shared weight matrix $W_{\text{out}} \in \mathbb{R}^{d \times d}$ to produce the final output embedding. This phase is standard and unchanged from a non-CARM attention implementation.

4.4 Concept embeddings

4.4.1 Phase 1.1: synthetic clustered embeddings

The initial implementation uses synthetic embeddings designed to occupy distinct subspaces by domain, mimicking the clustering structure of real semantic embeddings:

- Medical concepts (domain 1, TOSID prefix `0x10C5****`): strong signal in dimensions 0–20.
- Financial concepts (domain 2, TOSID prefix `0x11B1****`): strong signal in dimensions 21–41.
- PHI concepts (domain 3, TOSID prefix `0x13D2****`): strong signal in dimensions 42–63.

All embeddings are unit-normalised. This construction ensures that the routing policy and the embedding geometry are consistent — a deliberate choice for the initial validation, to confirm that routing correctly excludes concepts that would otherwise be semantically relevant to a given query.

4.4.2 Phase 1.2: GloVe embeddings at scale

Phase 1.2 of the implementation programme extended the system to 1,000 concepts with 100-dimensional GloVe embeddings [5], loaded from a pre-built binary format at startup. The vocabulary comprises 200 medical concepts, 200 financial concepts, 100 PHI concepts, and 500 neutral concepts. The number of attention heads was increased to 10 (head dimension 10).

This extension validates that the routing mechanism is not an artefact of synthetic embedding geometry. With real semantic embeddings, concepts from different domains are not cleanly separated in embedding space: a query about patient care will have non-trivial cosine similarity to financial concepts. The routing table must enforce access policy even when the model’s internal geometry would naturally attend to a blocked concept.

The $O(1)$ routing lookup was introduced in this phase, replacing the $O(n)$ linear scan used in Phase 1.1. The routing table is built once from rules at startup and held in memory for the duration of the session.

4.5 Test concept database

The test database used across both phases uses TOSID identifiers that illustrate the full identifier structure:

Concept	TOSID (hex)	Full TOSID string	Policy
aspirin	0x10C50001	10C-MED-SUP-ANB:PNC-AMP-500	ALLOW
ibuprofen	0x10C50002	10C-MED-SUP-ANL:IBU-400-TAB	ALLOW
morphine	0x10C50003	10C-MED-SUP-OPI:MOR-SUL-INJ	ALLOW
headache	0x10C50101	10C-MED-SYM-PAI:HEA-TEN-MIG	ALLOW
fever	0x10C50102	10C-MED-SYM-INF:PYR-ELE-TMP	ALLOW
billing_code	0x11B10001	11B-FIN-BIL-COD:ICD-10-DXX	DENY
insurance_claim	0x11B10002	11B-FIN-BIL-CLM:HSP-INS-FRM	DENY
payment_info	0x11B10003	11B-FIN-BIL-PAY:CRD-INF-DAT	DENY
patient_ssn	0x13D20001	11B-PER-IDN-GOV:USA-SSN-NUM	DENY
patient_name	0x13D20002	11B-PER-IDN-NAM:PAT-REC-FUL	DENY

The clinical context routing policy allows all concepts with TOSID prefix `0x10C5****` (medical domain) and denies all others. A single domain-level rule covers all medical concept variants; no per-concept rules are required.

4.6 Known issues and pre-publication fixes required

Two issues in the current codebase are noted for transparency and will be resolved prior to public release of the repository:

Macro redefinition. The `TOSID_GET_DOMAIN` macro is defined with different bit-shift values in `routing_table.h` (shift 24) and `multihead_attention.h` (shift 16). The compiler warns but substitutes correctly in each translation unit. The definitions will be unified in a single shared header.

Strncpy truncation warnings. Two `strncpy` calls copy exactly 63 bytes from 63-byte source strings. These are safe but trigger `-Wstringop-truncation` on GCC 13. They will be replaced with `memcpy` + explicit null termination.

5 Results

5.1 Experimental setup

All measurements were taken on a Linux x86-64 host running Ubuntu 24.04, compiled with `gcc 13.3.0 -O2 -march=native`. Timing uses `clock_gettime(CLOCK_MONOTONIC)`, which has a resolution of approximately 1 ns on this hardware. Each benchmark applies 1,000 warmup iterations before measurement to ensure instruction and data caches are in steady state. Mean values are reported over 10,000 measurement iterations unless otherwise stated.

A note on precision: routing latency values in the range 20–50 ns approach the resolution of the timer and should be interpreted as order-of-magnitude bounds rather than precise figures. All claims about routing overhead are stated conservatively, using the worst observed value rather than the mean.

5.2 Correctness: zero information leakage

The primary correctness property of CARM is C1: blocked concepts must receive exactly zero attention weight in every head, in every pass, regardless of query content.

To verify this, we ran 1,000 independent trials, each with a freshly drawn random query embedding (Box–Muller sampling, different seed per trial). The concept database contains 10 concepts: 5 accessible (medical domain) and 5 blocked (financial and PHI domains). In each trial we recorded the maximum attention weight assigned to any blocked concept, summed across all 8 heads.

Result: across all 1,000 trials — 5,000 individual blocked concept measurements, 8,000 individual head computations — the maximum attention weight on any blocked concept was **0.0000000000** (reported to 10 decimal places). No violation of C1 was observed.

Figure 2 shows the minimum attention weight on accessible concepts across all trials, confirming that the softmax distribution over the accessible set is well-behaved (weights in the range 0.188–0.200 for 5 accessible concepts, consistent with near-uniform distribution). The blocked concept weight, identically zero in every trial, is indicated by the dashed reference line at $y = 0$.




Figure 2: Zero-leakage verification over 1,000 trials with random queries. The plotted values are the minimum attention weight on accessible concepts. The dashed line at $y = 0$ represents the weight on blocked concepts, which was exactly zero in every trial. The gap between the two is absolute: no partial leakage was observed.

The zero result is structural, not statistical. Blocked concepts are absent from the softmax domain (Listing 4): they have no entry in the score array, so no floating-point rounding or numerical approximation can assign them nonzero weight. The 1,000-trial sweep confirms correct implementation of this structural property across varied query distributions.

5.3 Performance: routing overhead

Figure 3 shows routing overhead as a percentage of total computation time across concept counts from 5 to 100.

Routing overhead is below 0.13% across all tested concept counts, and decreases as the vocabulary grows. This behaviour is expected: attention time scales as $O(n \cdot d)$ in the number of accessible concepts, while routing time scales as $O(n)$ in vocabulary size with $O(1)$ per-concept lookup.

The worst-case overhead of 0.126% occurs at 5 concepts, where attention time is smallest ($24.5 \mu\text{s}$). At 100 concepts the overhead falls to 0.043%.

Figure 4 shows absolute routing and attention latencies. Routing latency ranges from 31 ns (5 concepts) to 159 ns (100 concepts), against attention times of $24.5\text{--}366 \mu\text{s}$: approximately three orders of magnitude apart across the tested range.

5.4 Latency stability

To assess timing stability, we collected 20,000 consecutive routing latency samples over the 10-concept clinical context. Table 1 reports the distribution.




Figure 3: Routing overhead as a percentage of total computation time (routing + attention), across concept counts 5–100. The dashed line marks the worst observed value of 0.126%, at 5 concepts. Overhead decreases as concept count grows because attention time scales with n while routing time scales sub-linearly.




Figure 4: Routing latency ($O(1)$ array lookup, solid) and linear scan baseline (dashed), both in nanoseconds, vs concept count. At this scale both methods are fast enough that differences fall within timer noise. The architectural advantage of $O(1)$ routing is confirmed at 1,000 concepts in Phase 1.2 of the programme (0.88 ns per lookup vs $23\times$ overhead for linear scan), not shown here.

The interquartile range of 3 ns (30–33 ns) indicates highly stable behaviour under normal OS scheduling. The p_{99} of 54 ns is the appropriate figure for latency budget planning; tail values reflect OS preemption rather than routing algorithm behaviour.

5.5 Policy effect on output

A routing policy that blocks concepts must produce a measurably different output embedding — otherwise the routing would be semantically inert. Figure 5 shows cosine distance between the output embedding under a blocking policy and under an unrestricted policy, as a function of the fraction of concepts blocked.

The near-zero distances at low blocking fractions (0.000030 at 10% blocked) reflect a geometric property of small vocabularies with random weights: removing one concept redistributes attention nearly uniformly, leaving the output close to the unrestricted case. This is not a failure of routing — C1 holds exactly — but a property of the unstructured embeddings in this proof-of-concept. In a trained model with structured semantic embeddings, blocking a relevant concept would produce substantially larger divergence.

	p_1	p_5	p_{10}	p_{25}	p_{50}	p_{75}	p_{90}	p_{95}	p_{99}	$p_{99.9}$	Mean
ns	30	30	30	30	31	33	41	46	54	182	35.8

Table 1: Routing latency percentiles over 20,000 samples, 10-concept clinical context. The interquartile range is 3 ns. The $p_{99.9}$ of 182 ns and maximum of 21,989 ns reflect OS scheduling interruptions, not algorithmic variance.




Figure 5: Cosine distance between the output embedding under a partial blocking policy and under an unrestricted policy, as a function of the percentage of concepts blocked (10-concept vocabulary). Distance reaches 1.0 at 100% blocked, where the output is the zero vector. The non-linearity around 60–90% reflects softmax redistribution geometry.

The 100% blocking case (distance 1.0) is exact by construction: no accessible concepts means a zero output vector, handled by graceful early exit (verified in Test 6a of the test suite).

5.6 Summary

Property	Claim	Observed	
Zero leakage (C1)	Blocked weight = 0 exactly	0.0000000000 across 5,000 measurements	✓
Routing overhead	< 1% at worst case	0.126% at worst	✓
Latency stability	$p_{99} < 100$ ns	$p_{99} = 54$ ns	✓
Policy effect	Measurable output divergence	Monotonically increasing with blocking	✓
Edge case: all blocked	Graceful handling	Early exit, zero vector	✓

Table 2: Summary of empirical results against stated claims.

6 Discussion and Ongoing Work

6.1 What has been established

The results of section 5 establish CARM as a functioning mechanism at the attention-component level. Three properties are confirmed.

The zero-weight guarantee is exact. Blocked concepts receive identically zero attention weight, not approximately zero. This follows from excluding them from the softmax domain before any score computation, and is confirmed across 5,000 independent measurements with varied

query distributions. The guarantee is structural: it holds for any query, including adversarially constructed ones, because the routing phase does not receive the query as input (C2).

The performance cost is negligible. Routing overhead of below 0.13% at worst case does not impose a meaningful constraint on system design. A mechanism with this overhead profile can be applied at every attention layer of a transformer stack without accumulating significant latency.

The mechanism is policy-driven, not model-driven. The same ACI component produces measurably different outputs under different AXI policies, without any change to its weights, and policy can be switched between passes in microseconds. This is the property that makes CARM practically useful in multi-tenant and multi-role deployments.

6.2 Identifier isolation and semantic isolation

A distinction central to the correct interpretation of CARM’s guarantees must be made explicit.

CARM provides identifier isolation. A concept identified by its TOSID is excised from the computation graph. Its attention weight is exactly zero. Its value vector contributes nothing to the output embedding. This guarantee is absolute within a CARM-compliant attention layer and is what the results of section 5.2 establish.

CARM does not claim to provide semantic isolation. When a concept is blocked, the softmax redistributes probability mass across the remaining accessible concepts. If the accessible vocabulary contains concepts that are semantically proximate to the blocked concept, the output embedding may retain indirect similarity to what it would have been had the blocked concept been accessible. This is not a violation of CARM’s guarantee — the blocked concept contributed nothing — but it is a real phenomenon with policy implications.

The boundary between these two properties is a design boundary, not a deficiency. Identifier isolation is CARM’s jurisdiction. Semantic isolation — ensuring that the accessible vocabulary does not contain proxies for blocked domains — is a policy design problem addressed at the level of the TOSID taxonomy. A well-designed TOSID hierarchy encodes semantic proximity explicitly, allowing an AXI policy to block not just a specific concept but its semantic neighbourhood. The two systems are complementary: CARM enforces the policy that TOSID makes expressible.

This separation of concerns is deliberate. A mechanism that attempted to enforce semantic isolation within the attention layer would require the routing decision to depend on inter-concept semantic relationships — which would violate C2 (query independence) and C3 (policy externality). CARM’s guarantees are achievable precisely because it operates on identifiers, not on meaning.

6.3 Limitations of the current implementation

The proof-of-concept operates over an explicit concept vocabulary with assigned TOSID identifiers. This differs from the operational conditions of a production system in several respects, each representing a named open problem in the research programme.

Array compaction vs. additive masking. The current implementation achieves the zero-weight guarantee by compacting the accessible array before the softmax — blocked concepts are absent from the computation rather than masked within it. This is mathematically equivalent to applying $-\infty$ additive masks to blocked logits in a static-geometry tensor, but operationally distinct: array compaction produces dynamically sized arrays per pass, while additive masking preserves static tensor geometry. The latter is required for dense matrix operations on GPU and TPU hardware. The additive masking variant is the subject of the next implementation phase and is expected to preserve the C1–C4 guarantees while enabling hardware acceleration.

Routing granularity. The $O(1)$ routing array operates on the upper 16 bits of the TOSID (domain + category prefix), which means policy is applied at the category level. All concept variants within a domain-category pair receive the same routing decision. Instance-level routing

— blocking a specific concept while allowing others in the same category — requires either a two-stage lookup or a different index structure. This is an architectural constraint, not an oversight: category-level policy is the appropriate granularity for most deployment scenarios, and instance-level exceptions are handled by subdividing categories in the TOSID hierarchy. The constraint is formalised here so that implementations requiring finer granularity know what to extend.

Concept representation in trained models. In a trained LLM, meaning is distributed across token embeddings and emerges from the interaction of layers. There is no discrete vocabulary of concepts with explicit identifiers; there are tokens, and concepts emerge from their combinations. Applying CARM to a full LLM requires resolving the mapping between TOSID-identified concepts and the token-level representations the model uses. Two approaches are under investigation: annotating token vocabularies with TOSID prefixes at the embedding level, and introducing a lightweight concept extraction layer that maps token contexts to TOSID-identified concept activations prior to the attention computation.

Residual and feed-forward paths. The current implementation governs only the attention sublayer. A transformer block also contains a feed-forward sublayer and residual connections. Whether blocked concepts can re-enter the computation through these paths is an empirical question that requires full-stack integration to answer. The architecture does not assume this cannot happen; it assumes it must be measured, and if necessary addressed by applying CARM at additional points in the computation graph.

Autoregressive generation dynamics. In sequential token generation, each step’s output becomes part of the context for the next step. A concept blocked at step t may influence the context embedding at step $t + 1$ through tokens already generated. Whether this constitutes a meaningful policy bypass or merely an indirect semantic influence is a question sequential generation experiments can answer.

Vocabulary scale. The proof-of-concept operates over vocabularies of 10–1,000 concepts. Production systems may require vocabularies of tens of thousands of TOSID-identified concepts. The $O(1)$ routing mechanism scales to this range without algorithmic change, but the policy management problem — how to assign, maintain, and version TOSID identifiers across a large vocabulary — requires tooling beyond the scope of this paper.

6.4 The feedback path and TOSIDcoders

The ACI/AXI architecture as presented in section 3 is unidirectional within a single attention pass: policy flows from AXI through CARM to ACI, with no feedback path. This is by design; the unidirectional flow is what makes the C1–C4 guarantees achievable.

Across passes, however, a complete system requires a feedback path: the AXI component must be able to observe what the ACI component attended to, reason about it in TOSID terms, and update routing policy for subsequent passes. This requires translating from the continuous, distributed output space of the ACI component back into the discrete, hierarchical, policy-addressable space of TOSID — a non-trivial mapping problem.

We identify this as requiring a purpose-trained class of model, which we term a **TOSIDcoder**. The analogy to the role of embedding models in retrieval-augmented generation (RAG) [4] is precise: just as RAG required a separate model class trained specifically to map text into a metric space suitable for retrieval, a CARM/TOSID system with feedback requires a model trained specifically to map ACI outputs into TOSID concept space. The generator (ACI) and the grounder (TOSIDcoder) have different training objectives and should be separate systems.

A TOSIDcoder operating at the ACI output boundary would allow the AXI component to ask: which TOSID-identified concepts does this output activate? Does the output carry semantic content from blocked domains, even indirectly? Should routing policy be tightened for the next pass? These are answerable questions once ACI outputs are grounded in TOSID space. They

are not answerable with the current architecture, which is why the feedback path is named here as an open component rather than a solved one.

TOSIDcoders also resolve the semantic isolation question of section 6.2. The question “did a blocked concept leak semantically through proxy concepts?” becomes measurable: run the output through a TOSIDcoder and inspect the resulting concept distribution. If blocked-domain concepts appear above threshold, the policy has been semantically circumvented and the AXI component can respond — without CARM’s guarantees having been violated, since identifier isolation held throughout.

6.5 The research programme

The work reported in this paper is the first in a sequence of planned investigations. The programme is structured around four named stages, each building on the previous and each producing a publishable contribution.

Stage I (this work): Mechanism definition and component-level proof. CARM is formally defined, implemented in C over a multi-head attention layer, and its correctness and performance properties are established at the component level. The ACI/AXI architectural model is introduced. The proof-of-concept uses array compaction; the additive masking variant is deferred to Stage II.

Stage II: Hardware-compatible implementation and full-stack contact. Two parallel investigations. First, the array compaction implementation is replaced by an additive $-\infty$ masking variant operating over static tensor geometry, making the mechanism compatible with GPU and TPU dense matrix kernels. The C1–C4 guarantees are re-verified under this variant and timing metrics are re-established. Second, CARM is applied to a complete transformer block — residual connections, feed-forward sublayer, layer normalisation — and the fate of the zero-weight guarantee through the full computation graph is measured empirically. Sequential generation with context accumulation is tested, along with adversarial robustness: can the routing policy be bypassed through prompt construction, semantic proximity, or multi-step reasoning? The honest result of Stage II is either (a) the guarantees survive full-stack contact, in which case CARM is retrofittable to existing architectures, or (b) the guarantees are compromised by structural properties of existing transformer architectures, in which case Stage III is warranted.

Stage III: Architecture. If Stage II establishes that existing transformer architectures structurally cannot accommodate CARM’s guarantees without compromise, Stage III proposes an architecture in which CARM is a first-class primitive rather than a retrofit. In such an architecture, the ACI/AXI separation is built in from the ground up: attention layers are designed around the routing interface, residual paths are structured to preserve routing decisions, and the feed-forward sublayer is either CARM-aware or replaced by an alternative. This stage may draw on the PTAC (Printed Toroidal Array Computer) programme for hardware-level TCAM integration, which provides a natural physical implementation of the $O(1)$ routing table. Stage III is not contingent on Stage II finding failure; it may proceed in parallel if Stage II results suggest that retrofitting, while possible, imposes unacceptable architectural compromises.

Stage IV: TOSIDcoders and the feedback loop. The final stage of the current programme closes the ACI/AXI feedback path by developing the TOSIDcoder model class. A TOSIDcoder is trained to map from ACI output space into TOSID concept space, enabling the AXI component to reason about ACI outputs in policy-addressable terms and update routing for subsequent passes. Training objectives, corpus requirements, bootstrapping strategies, and evaluation metrics for TOSIDcoders are the primary research questions. Stage IV presupposes a stable TOSID taxonomy (developed in parallel under the TOSID programme) and a validated CARM implementation from Stages I–III. The completion of Stage IV constitutes a fully operational closed-loop ACI/AXI system.

6.6 Relation to broader AI safety work

CARM addresses a specific, narrow problem: enforcing categorical concept-level access policy at the attention layer. It is not a general AI safety mechanism and does not claim to be. It does not prevent a model from producing harmful outputs through reasoning that does not involve the blocked concepts; it does not address deceptive alignment, reward hacking, or distributional shift. Its scope is precisely its claim: within a CARM-compliant attention layer, no denied concept contributes to output.

Within that scope, however, the guarantee is qualitatively different from what behavioural safety mechanisms can offer. RLHF [2], constitutional AI [3], and output filtering produce probabilistic tendencies; CARM produces a structural constraint. For applications where categorical policy compliance is required — regulated industries, information barriers, jurisdictional constraints — structural constraints are the appropriate tool, and the present work demonstrates that they are achievable at acceptable performance cost.

The research programme described in section 6.5 is not an attempt to retrofit CARM onto existing transformer architectures as they stand. It is an investigation into whether those architectures are compatible with the kind of structural guarantees that regulated deployment requires. The answer may be yes, or it may be no. If no, this is not a finding against CARM but a finding against the suitability of current architectures for policy-compliant deployment — a result with significant implications for how AI systems in regulated domains should be built from the ground up.

6.7 Conclusion

We have introduced CARM as a formally defined mechanism (Definition 2.1, criteria C1–C4), situated it within the ACI/AXI architectural model, implemented it in C over a multi-head attention layer with real semantic embeddings, and measured its correctness and performance properties.

The central result is that architectural concept isolation is achievable with negligible performance cost. The routing overhead of below 0.13% at worst case, combined with exact zero-weight enforcement confirmed across 5,000 independent measurements, establishes CARM as a viable mechanism for deployment in attention-based systems where policy compliance must be guaranteed rather than approximated.

CARM is the first component of a broader research programme aimed at establishing what structural guarantees are achievable in AI systems, and what architecture is required to achieve them. The programme is active; the open questions are named; the path forward is clear.

The implementation is available at <https://github.com/ha1tch/carm>.

References

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in Neural Information Processing Systems* **30** (2017).
- [2] P. Christiano, J. Leike, T. B. Brown, M. Martic, S. Legg, and D. Amodei, “Deep reinforcement learning from human preferences,” *Advances in Neural Information Processing Systems* **30** (2017).
- [3] Y. Bai, S. Kadavath, S. Kundu, A. Askell, J. Kernion, A. Jones, A. Chen, A. Goldie, A. Mirhoseini, C. McKinnon, et al., “Constitutional AI: Harmlessness from AI feedback,” *arXiv preprint arXiv:2212.08073* (2022).
- [4] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” *Advances in Neural Information Processing Systems* **33** (2020).
- [5] J. Pennington, R. Socher, and C. Manning, “GloVe: Global vectors for word representation,” *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1532–1543, 2014.
- [6] P. Jackson, *Introduction to Expert Systems*. Addison-Wesley, 1986.